

Introduction to MongoDB

Storing Data

Nuri Halperin
@nurih



Overview


Storage Engine

Saving documents

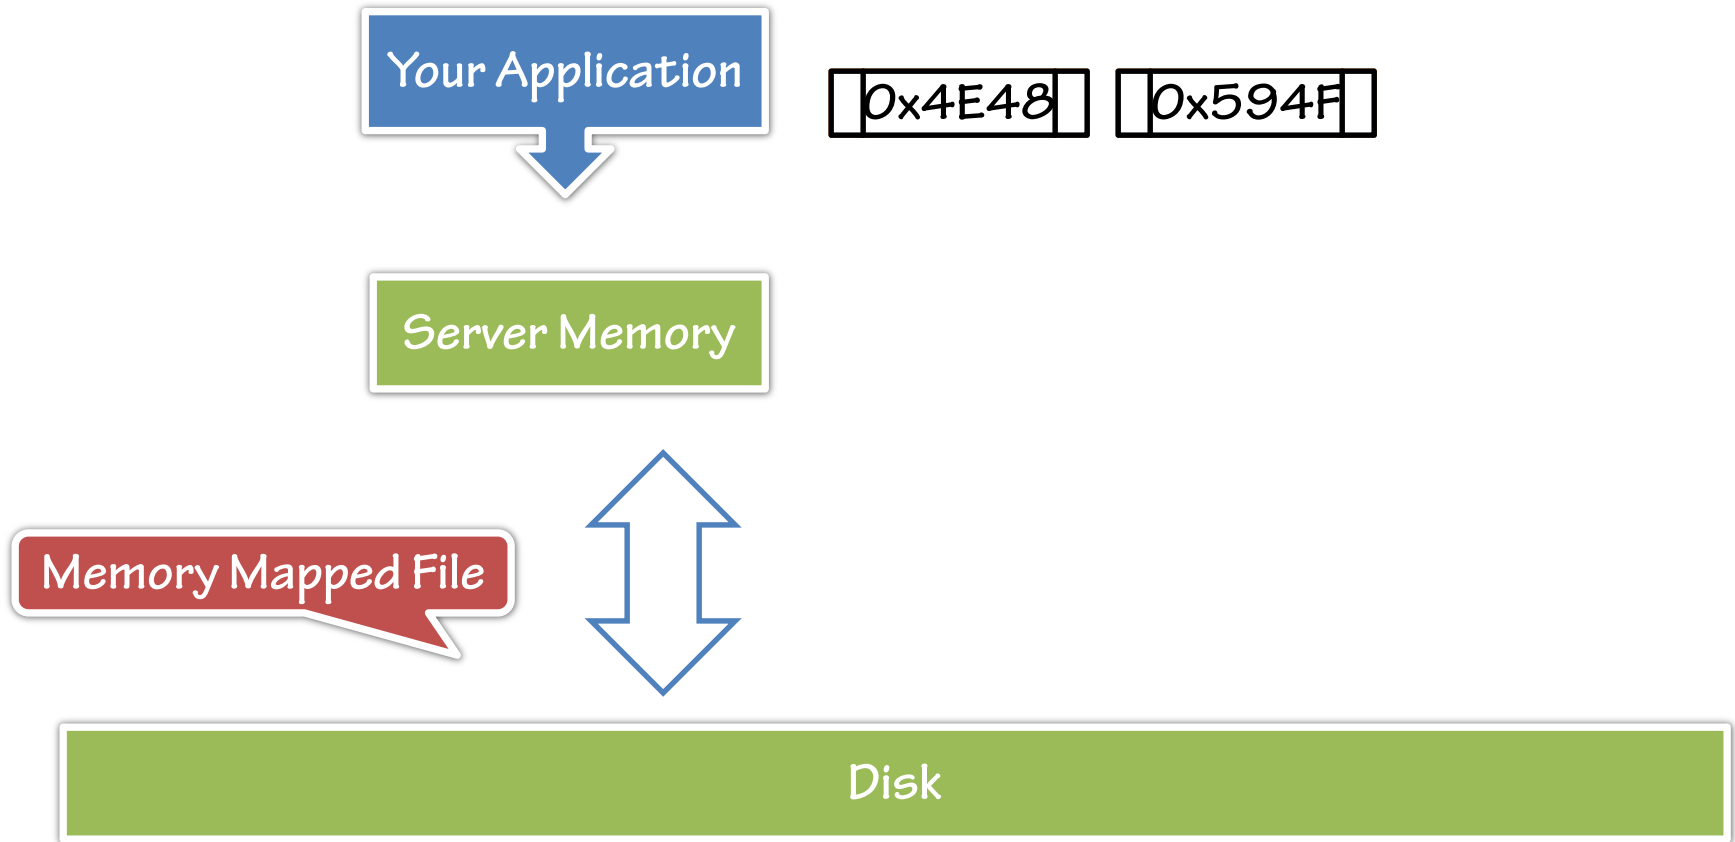
Updating documents

Storage Internals

- How is data stored?



Storage Internals



BSON { 01010100 11101011 10101110 01010101 }

BSON [*bee · sahn*], short for Binary [JSON](#), is a binary-encoded serialization of JSON-like documents. Like JSON, BSON supports the embedding of documents and arrays within other documents and arrays. BSON also contains extensions that allow representation of data types that are not part of the JSON spec. For example, BSON has a Date type and a BinData type.

BSON can be compared to binary interchange formats, like [Protocol Buffers](#). BSON is more "schema-less" than Protocol Buffers, which can give it an advantage in flexibility but also a slight disadvantage in space efficiency (BSON has overhead for field names within the serialized data).

BSON was designed to have the following three characteristics:

1. Lightweight

Keeping spatial overhead to a minimum is important for any data representation format, especially when used over the network.

2. Traversable

BSON is designed to be traversed easily. This is a vital property in its role as the primary data representation for [MongoDB](#).

3. Efficient

Encoding data to BSON and decoding from BSON can be performed very quickly in most languages due to the use of C data types.

[specification](#)[implementations](#)[FAQ](#)[discussion](#)

Saving Data

#1: A document must have an **_id** field.

#2: There are very few rules.

Update

Atomic within a document

Targets individual field(s)

Rich syntax support

Update

Which
Document?

```
db.foo.update(query, update, options);
```

Collection Name

What
Change?

One?
Many?
Upsert?

Find And Modify

Collection Name

Which Document?

```
db.foo.findAndModify( {  
  query: <document>,  
  update: <document>,  
  upsert: <boolean> ,  
  remove: <boolean>,  
  new: <boolean>,  
  sort: <document>,  
  fields: <document>} );
```

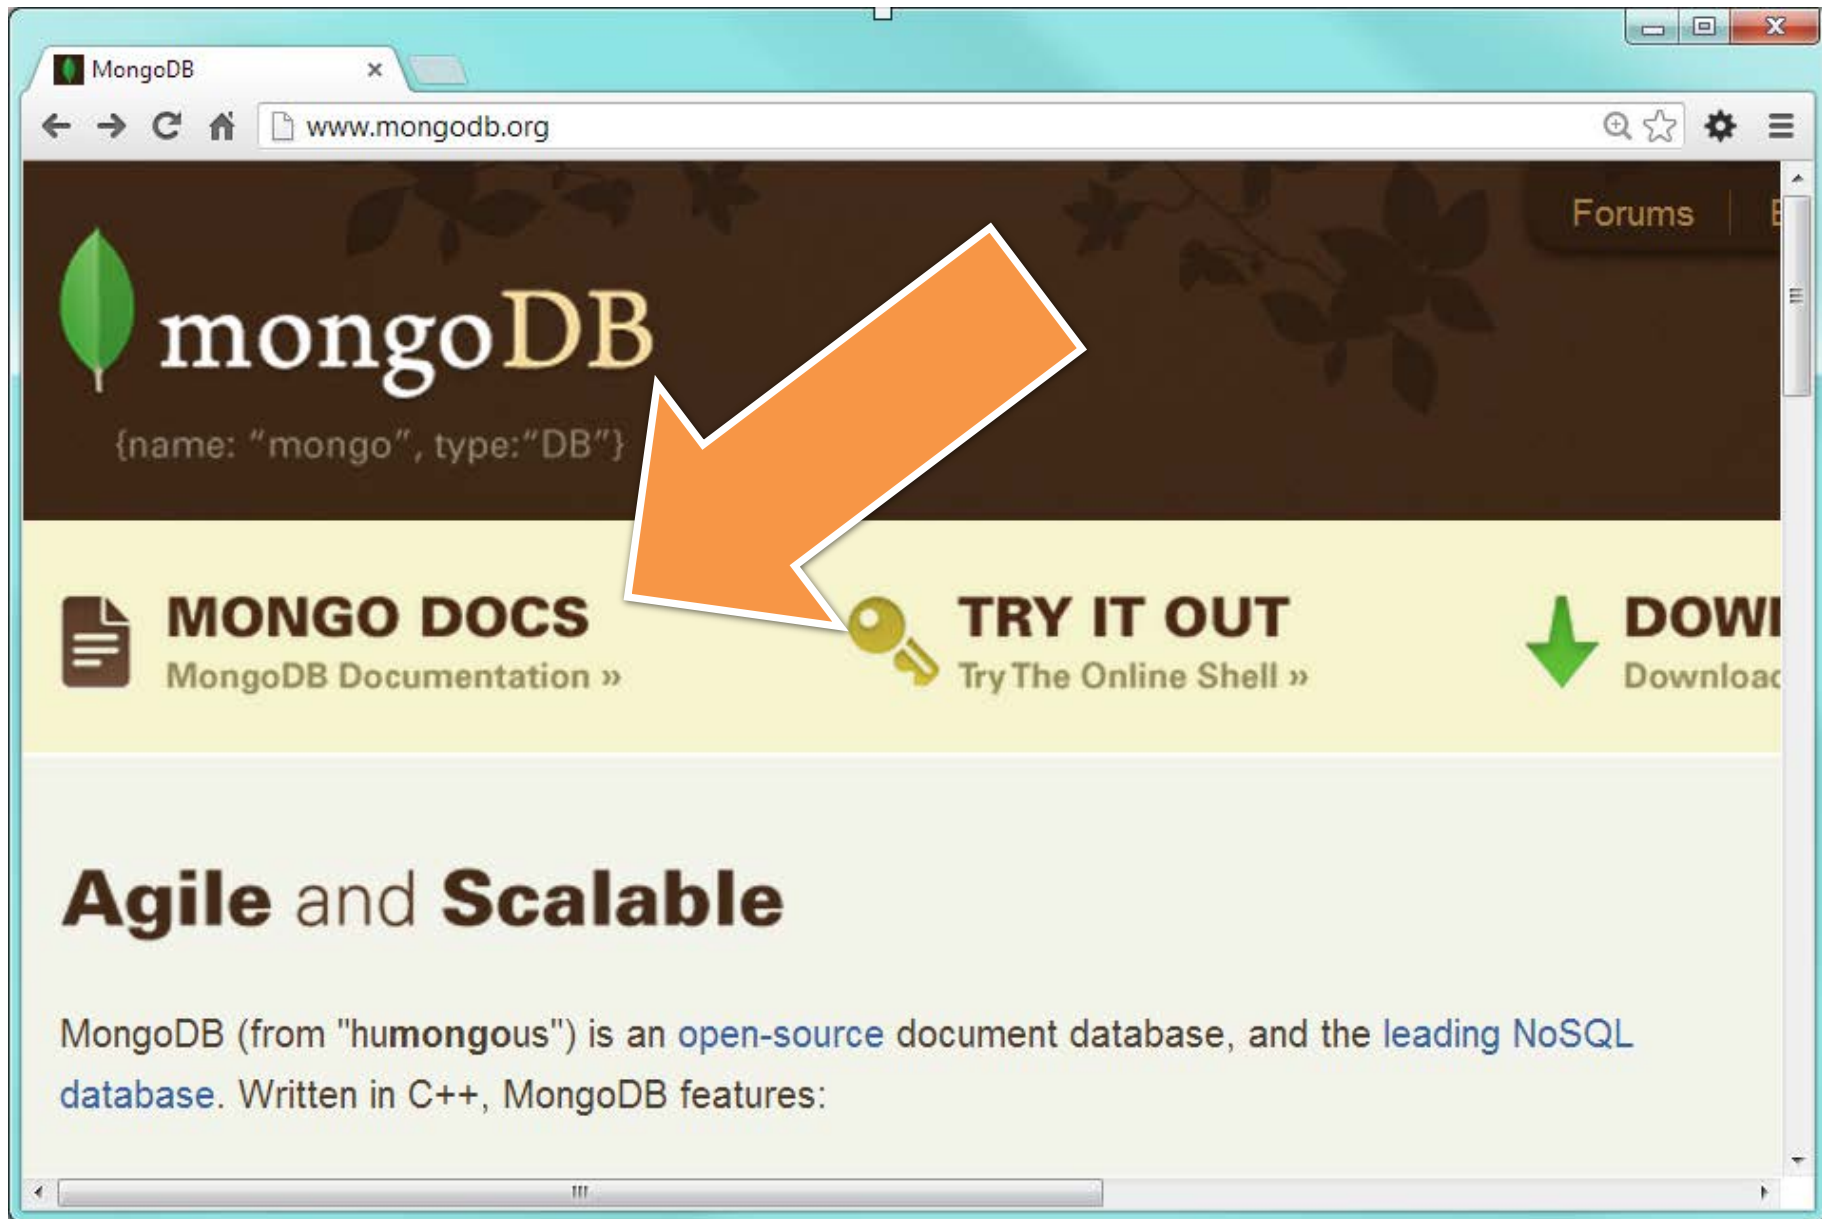
What Change?

Delete it?

Return new/old?

Query order

Return what?



Summary

Storage

- BSON
- Memory Mapped File

New Documents

- Save
- Insert
- Concurrency issue

Updates

- Update
- FindAndModify